

## Cover letter

JULY 14, 2026  
BRAINTRUST / PRINCIPAL ROLE

Hiring Team  
Braintrust and its AI-lab client  
United States - Remote

Dear Hiring Team,

The expensive mistake in AI data work is rarely a shortage of activity. It is allowing collection, annotation, or evaluation to scale before the team has agreed on what the model should do, which cases matter, who has judgment authority, how disagreement will be handled, and what evidence earns an acceptance decision. That is the problem I would help Braintrust and its client solve.

My career has repeatedly centered on making quality legible before execution becomes difficult to reverse. At DudeWorth, I design AI-readiness and agentic workflow patterns that combine structured context, retrieval, human judgment, escalation, evaluation, and governance. At Vape-Jet, I connect production, customer, support, engineering, and machine-fleet signals into earlier risk detection, standard work, training, and accountable feedback loops across approximately 1,000 fielded robots. Earlier, at Compunetics, I built quality systems, root-cause practices, performance measures, and process controls in high-reliability engineering and manufacturing. At Amazon, I built operating mechanisms for a 24/7 environment supporting approximately five million second-day orders annually.

The title history is not a conventional AI-research sequence, and I will not inflate it into one. What transfers directly is the operating pattern this advisory role requires: discover the real decision, expose failure modes early, define observable quality, establish human authority, calibrate the system, and give program managers a design that can scale without hiding uncertainty.

My independent candidate vision makes that reasoning visible through a Judgment Convergence Map, an Evaluation Design Contract, and a first-90-days plan. The core thesis is simple: **define the judgment before you scale the data.**

I would welcome the opportunity to test that thesis against a real project and show how I would turn an ambiguous goal into a decision-useful dataset, rubric, calibration plan, and pilot.

Carpe diem,  
**Russell Dudek**