Russell Dudek

Interview Thesis Brief

Braintrust

Independent candidate materials
Principal AI Data & Evaluation Strategist

Pittsburgh, Pennsylvania 412.287.8640 russelldudek@gmail.com linkedin.com/in/russelldudek russelldudek.github.io/braintrust

Principal AI Data & Evaluation Strategist

Define the judgment before you scale the data.

A project is ready to scale when the intended behavior, coverage, ground-truth authority, rubric, calibration process, and acceptance/stop rules form one coherent decision system.

Real mandate

Serve as the upstream quality architect: translate business and research goals into data requirements and evaluation systems before program execution begins.

Company moment

Braintrust has expanded from a talent marketplace into AI infrastructure, including human data for frontier AI work. The role strengthens the design layer between expert contributors, program management, and AI-lab stakeholders.

Central tensions

- Speed to pilot vs measurement validity
- Reusable methods vs domain specificity
- Ground truth vs legitimate ambiguity
- Expert judgment vs scalable throughput
- Advisory influence vs delivery authority

Strategic value

Prevent quality debt: the hidden rework, inconsistent judgments, unusable data, and weak decisions created when ambiguity is discovered after collection begins.

THE PRINCIPAL OPERATING STANCE

Before discovery ends	**Before sourcing begins**
Make the decision, target behavior, risk, stakeholders, and unknowns explicit.	Define domain expertise, data provenance, coverage, confidentiality, and evaluator qualification.
Before calibration	**Before pilot acceptance**
Operationalize the rubric and create anchored examples that distinguish adjacent scores.	Review validity, coverage, agreement, adjudication burden, decision utility, and unresolved risk.

Public context: Braintrust About, AI Training, and Responsible AI pages. Role details: supplied job description and posting URL.

Evaluation Design Contract

1. OUTCOME CLAIM

The decision the dataset/evaluation must support, the user of that decision, and the consequence of being wrong.

2. BEHAVIOR AND COVERAGE

Capabilities, failure modes, domains, difficulty, populations, languages, rare cases, and exclusions.

3. GROUND-TRUTH AUTHORITY

Separate factual, expert, preference, policy, and plural judgments; define who resolves each.

4. HUMAN-EVALUATION RUBRIC

Observable criteria, scoring anchors, examples, rationale requirements, and prohibited shortcuts.

5. CALIBRATION AND ADJUDICATION

Qualification, gold items, disagreement review, drift checks, feedback, and change control.

6. ACCEPTANCE AND STOP RULES

Decision criteria, confidence, coverage, escalation, narrowing, redesign, and termination conditions.

QUALITY SCORECARD

Validity

Does the measure represent the intended behavior?

Coverage

Are material slices and failure modes represented?

Judgment quality

Agreement, rationale, drift, and adjudication burden.

Decision utility

Does the result lead to an explicit model/data/product decision?

Economics

Expert time, rework, cycle time, and avoidable review load.

Learning velocity

How quickly pilot evidence improves the design.

Proof map, objection, and questions

DudeWorth

AI readiness, LLM/RAG/agentic workflow patterns, human review, evaluation, escalation, and governance.

Vape-Jet

Cross-system quality and feedback mechanisms joining human, customer, operational, and machine signals.

Compunetics

High-reliability quality systems, root cause, process controls, technical translation, and consequence-aware execution.

Amazon

Customer-backward operating mechanisms, standard work, visibility, escalation, and scaled execution.

STRONGEST OBJECTION

Concern: no verified claim of prior ownership of a frontier-lab evaluation program or 10+ formal years in AI. **Response:** evaluate the directly relevant work pattern: upstream definition of quality, technical-business translation, human judgment, risk, calibration, operating mechanisms, and AI workflow design. The first 90 days are structured to prove that transfer quickly.

EXECUTIVE INTERVIEW QUESTIONS

01

Where does ambiguity currently create the most rework?

02

Which model or data decisions must the evaluation support?

03

Who owns the definition of acceptable quality?

04

How are expert disagreements preserved and adjudicated?

05

What coverage gap would be most consequential?

06

Which quality methods should be reusable across programs?

07

What evidence persuades researchers, engineering, and executives?

08

How is rubric drift detected after launch?

09

When should a project narrow scope rather than collect more data?

10

What would outstanding performance look like after three engagements?